



Uni.lu HPC School 2021

PS3: Getting Started 2.0: SLURM, performance engineering and basic launchers

Uni.lu High Performance Computing (HPC) Team
H. Cartiaux

University of Luxembourg ([UL](#)), Luxembourg

<http://hpc.uni.lu>



Latest versions available on Github:



UL HPC tutorials:

<https://github.com/ULHPC/tutorials>

UL HPC School:

hpc.uni.lu/education/hpcschool

PS3 tutorial sources:

ulhpc-tutorials.rtfd.io/en/latest/beginners



Summary

1 Introduction

2 Getting Started on ULHPC

3 Batch Scheduling Configuration (Slurm)

Slurm commands

ULHPC Slurm Configuration

Usage Example

Basic Slurm Launchers

4 User [Software] Environment

Main Objectives of this Session

- Review ULHPC environment
- **Design and usage of SLURM**
 - ↪ global configuration and Aion specificities
 - ↪ **common** and **advanced** SLURM tools and commands
 - ✓ srun, sbatch, squeue etc.
 - ✓ job specification
- Understand User Software environment
 - ↪ LMod modules, generated with Easybuild
- SLURM generic **launchers** you can use for your own jobs



Summary

1 Introduction

2 Getting Started on ULHPC

3 Batch Scheduling Configuration (Slurm)

Slurm commands

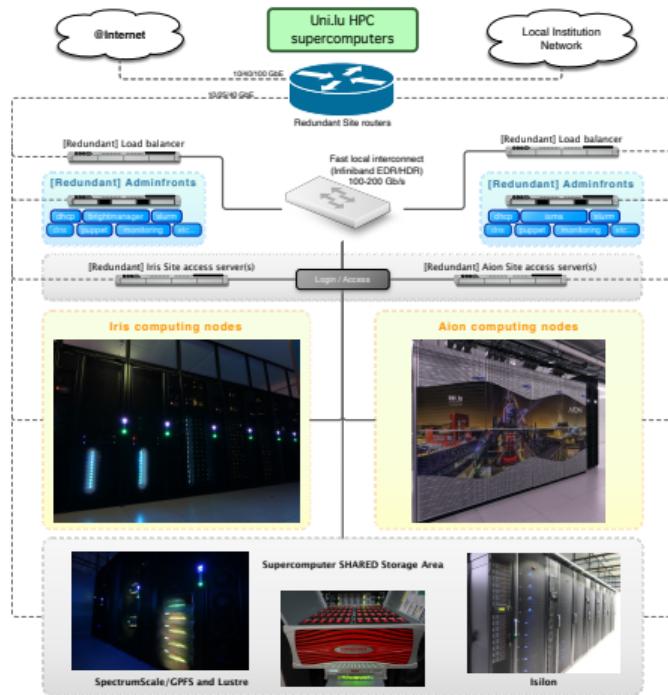
ULHPC Slurm Configuration

Usage Example

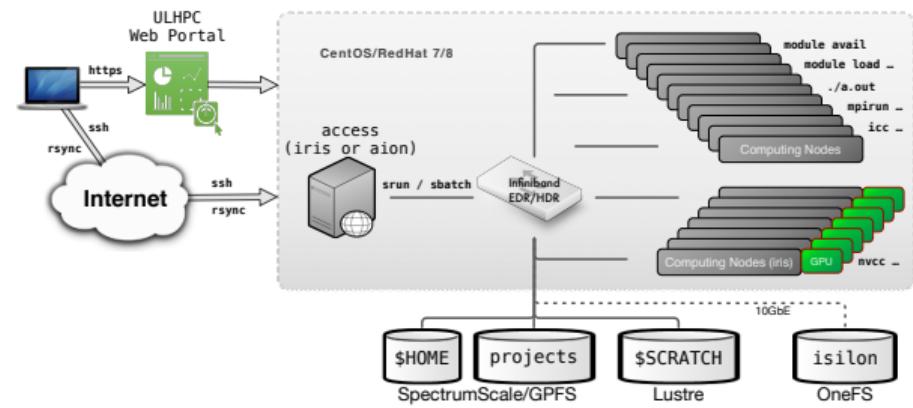
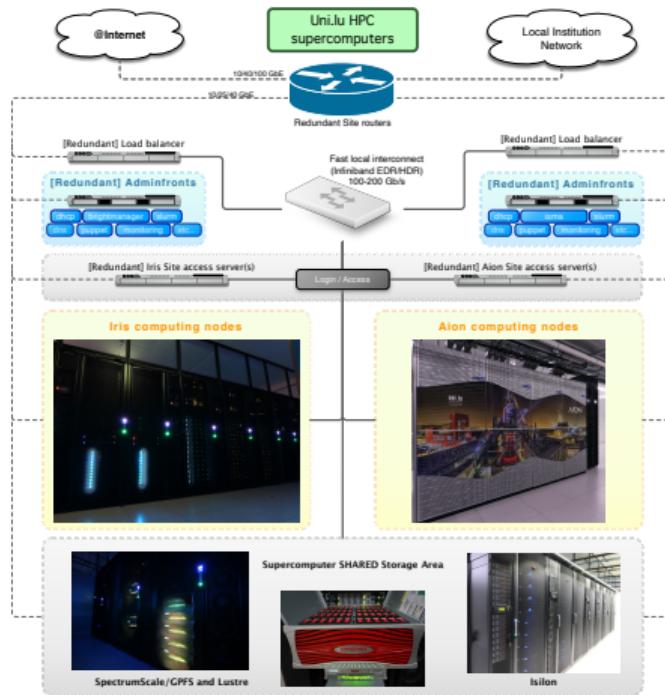
Basic Slurm Launchers

4 User [Software] Environment

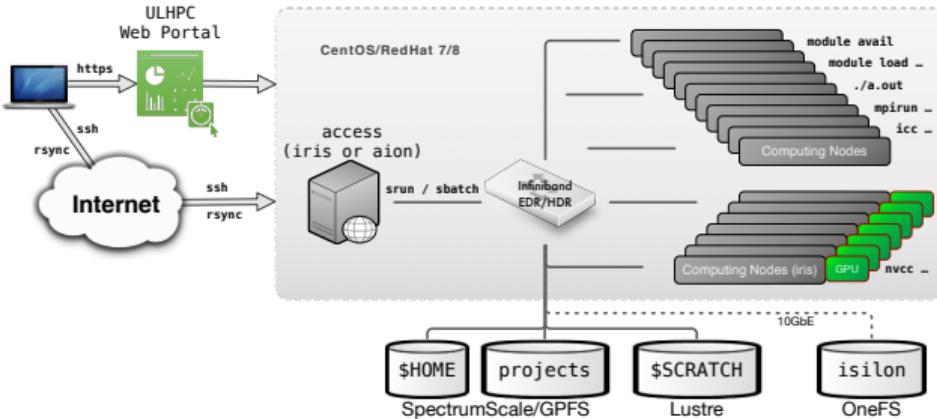
UL HPC Supercomputers: General Architecture



UL HPC Supercomputers: General Architecture



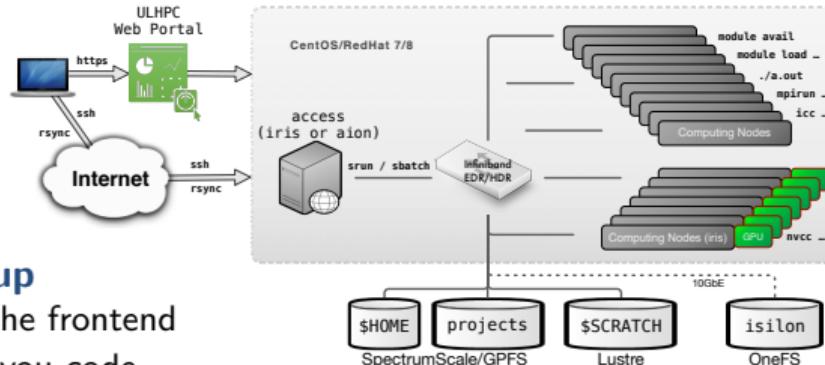
Compute Nodes / Storage Environment



- **Storage usage:** `df-ulhpc [-i]`
 - ↳ \$HOME: regular backup policy
 - ↳ \$SCRATCH **NO** backup & purged
 - ✓ 60 days retention policy
 - ↳ Project quotas attached to group
 - ✓ **not** (*default*) clusterusers group
 - ✓ Commands writing in project dir: `sg <group> -c "<command>"`
- **LMod/Environment modules**
 - ↳ **Not** on access, **only** on compute nodes

Directory	FileSystem	Max size	Max #files	Backup
\$HOME (iris)	GPFS	500 GB	1.000.000	YES
\$SCRATCH	Lustre	10 TB	1.000.000	NO
Project	GPFS	<i>per request</i>		PARTIALLY (/backup subdir)
Project	OneFS	<i>per request</i>		PARTIALLY

Typical Workflow on UL HPC resources

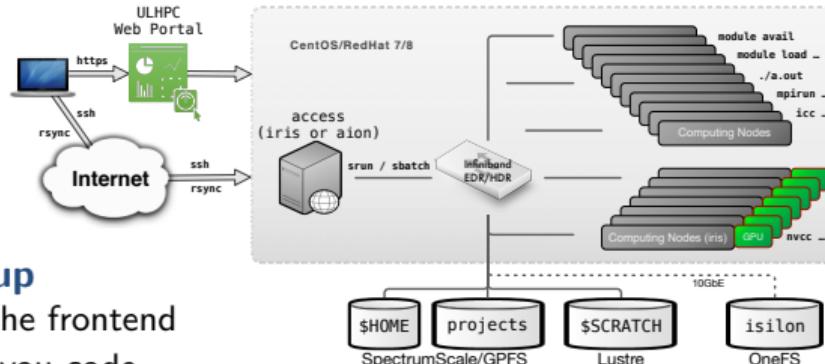


- Preliminary setup

- ① Connect to the frontend
- ② Synchronize your code
- ③ Reserve a few interactive resources
 - ✓ (eventually) build your program
 - ✓ Test on small size problem
 - ✓ Prepare a launcher script

ssh, screen
 scp/rsync/svn/git
 srun -p interactive [...]
 gcc/icc/mpicc/nvcc...
 srun/python/sh...
 <launcher>. {sh|py}

Typical Workflow on UL HPC resources



- Preliminary setup

- ① Connect to the frontend
- ② Synchronize your code
- ③ Reserve a few interactive resources
 - ✓ (eventually) build your program
 - ✓ Test on small size problem
 - ✓ Prepare a launcher script

- Real Experiment

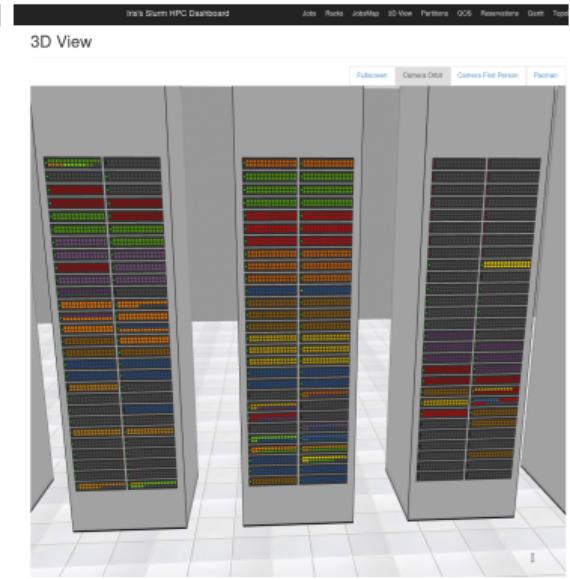
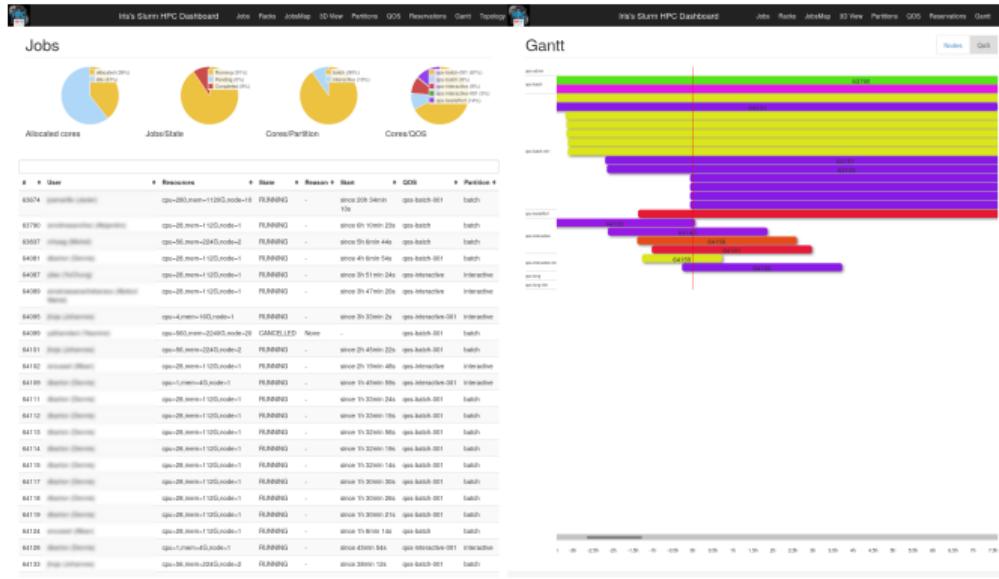
- ① Reserve passive resources
- ② Grab the results

ssh, screen
 scp/rsync/svn/git
 srun -p interactive [...]
 gcc/icc/mpicc/nvcc...
 srun/python/sh...
 <launcher>. {sh|py}

sbatch [...] <launcher>
 scp/rsync/svn/git ...

Slurm Web

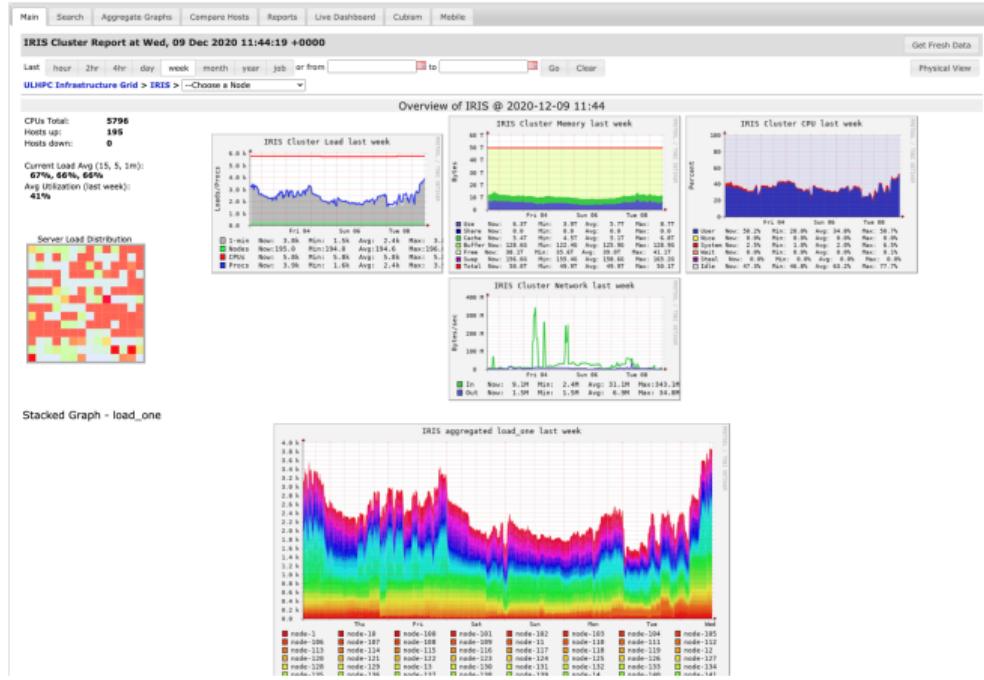
(internal network only)



Getting Started on ULHPC

Ganglia

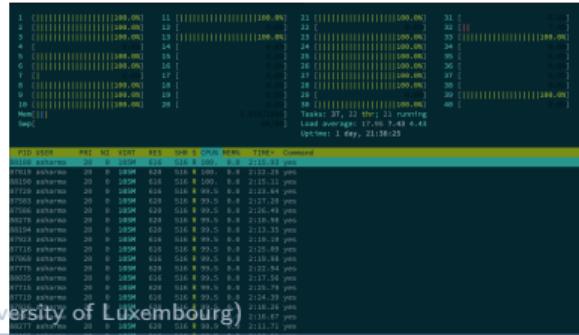
<https://hpc.uni.lu/iris/ganglia/>



- **WIP:** Prometheus / Graphana
→ together with aion

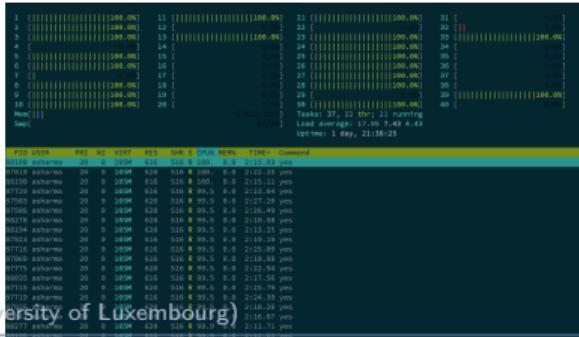
Interactive Monitoring/debugging

```
$> sq  
# Connect to your job (first allocated node)  
$> sjoin <JOBID>  
$> htop      # view of all processes  
              # F5: tree view  
              # u <name>: filter by process of <name>  
              # q: quit
```



Interactive Monitoring/debugging

```
$> sq
# Connect to your job (first allocated node)
$> sjoin <JOBID>
$> htop # view of all processes
          # F5: tree view
          # u <name>: filter by process of <name>
          # q: quit
```



• Arm Forge suit

→ ddt, map, perf-report

ARM REPORTS

Command: /home/luhpcusers/lehnhausen/assembly/arm-mesh/arm-mesh-a-setor_small.tst
Resources: 1 node (28 physical, 28 logical cores per node)
Memory: 28 processes, 28 threads per node
Machine: iis.apr.7.2020.13:14:20 (RTC#92)
Start time: 2020-04-07T13:14:20Z
Full path: /home/luhpcusers/lehnhausen/BPG/INTEL_PERFORMANCE/ARM/Carde...



Summary: heart_demo is MPI-bound in this configuration

Compute	47.3%	
MPI	52.6%	
I/O	<8.1%	

This application ran was MPI-bound. A breakdown of this time and advice for investigating further is in the MPI section below.

CPU

A breakdown of the 47.3% CPU time:

Single-core code: 37.0%

OpenMP regions: 2.1%

Scalar numeric ops: 5.1%

Vector numeric ops: 0.2%

Memory accesses: 74.6%

Per-process performance can be analyzed by serial sections of the application. Use a profiler to find these or run with fewer threads and more processes.

The per-core performance is memory-bound. Use a profiler to identify memory-consuming loops and check their cache performance.

MPI

A breakdown of the 52.6% MPI time:

Time in collective calls: 28.7%

Time in point-to-point calls: 71.3%

Effective process point-to-point rate: 2.0 Gbps

Most of the time is spent in point-to-point calls with a low transfer rate. This can be caused by inefficient message sizes, such as many small messages, or by imbalanced workloads causing processes to wait.

I/O

A breakdown of the <8.1% I/O time:

Time in reads: 0.0%

Time in writes: 0.0%

Effective process read rate: 0.00 bytes/s

Effective process write rate: 0.00 bytes/s

Most of the time is spent in write operations with a very low rate. This can be caused by inefficient file operations or the filesystem or inefficient access patterns. Use an I/O profiler to investigate which write calls are inefficient.

OpenMP

A breakdown of the 2.1% time in OpenMP regions:

Computation: 0.1%

Synchronization: 99.9%

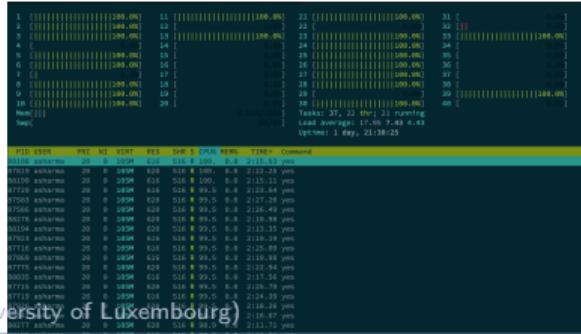
Physical thread utilization: 100.0%

System load: 355.1%

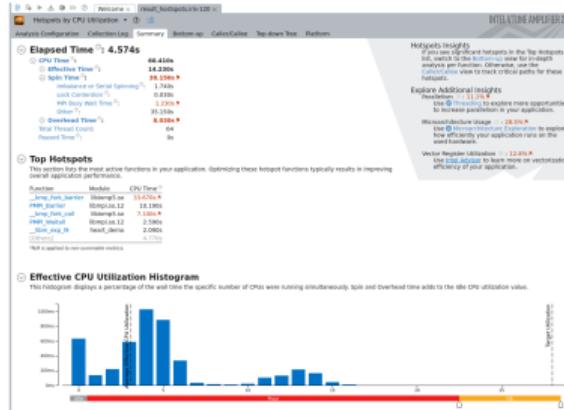
The system load is high. Please be kind to system resources.

Interactive Monitoring/debugging

```
$> sq  
# Connect to your job (first allocated node)  
$> sjoin <JOBID>  
$> htop # view of all processes  
        # F5: tree view  
        # u <name>: filter by process of <name>  
        # q: quit
```



- Arm Forge suit
 ↪ ddt, map, perf-report
 - Intel VTune/Advisor/Inspector



Discovering and visualizing UL HPC environment

Your Turn!

Hands-on Discovering your HPC environment

▶ url ◀ | [github](#) | [src](#)

- Connect to the cluster frontend
- Navigate between \$HOME, \$SCRATCH
- Discovering storage usage
- Live status tools (Slurm Web, Ganglia...)

Summary

1 Introduction

2 Getting Started on ULHPC

3 Batch Scheduling Configuration (Slurm)

Slurm commands

ULHPC Slurm Configuration

Usage Example

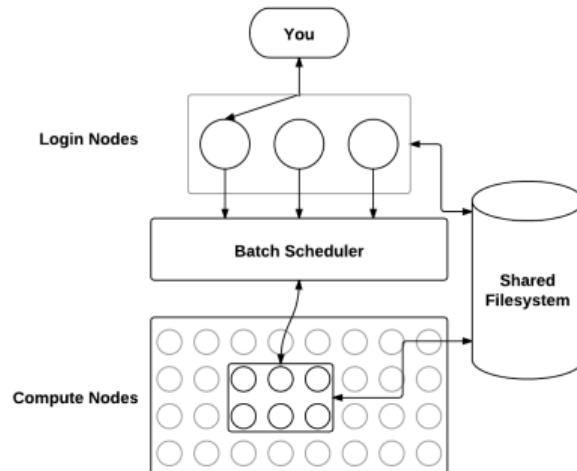
Basic Slurm Launchers

4 User [Software] Environment

Resource and Job Management Systems

- **Resource and Job Management System (RJMS)**

- ↪ *Glue* for a parallel computer to execute parallel jobs
- ↪ **Goal:** satisfy users demands for computation
 - ✓ assign resources to user jobs with an efficient manner



Resource and Job Management Systems

- **Resource and Job Management System (RJMS)**

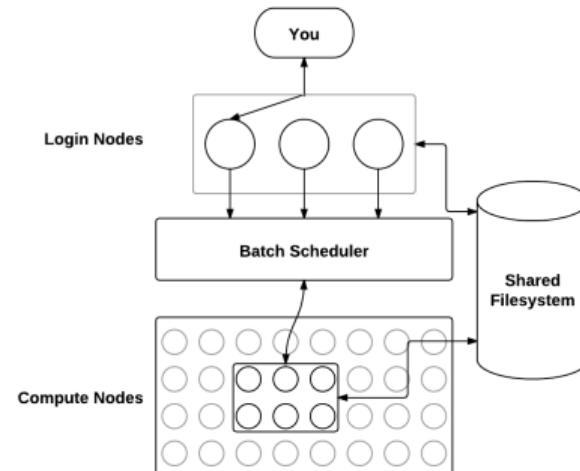
- ↪ *Glue* for a parallel computer to execute parallel jobs
- ↪ **Goal:** satisfy users demands for computation
 - ✓ assign resources to user jobs with an efficient manner

- **HPC Resources:**

- ↪ Nodes (typically a unique IP address)
 - ✓ Sockets / Cores / Hyperthreads
 - ✓ Memory
 - ✓ Interconnect/switch resources
- ↪ Generic resources (e.g. GPUs)
- ↪ Licenses

- **Strategic Position**

- ↪ Direct/constant knowledge of resources
- ↪ Launch and otherwise manage jobs



Slurm on ULHPC clusters

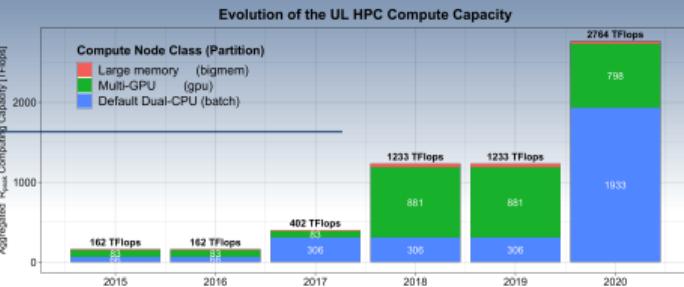
- ULHPC uses **Slurm** for cluster/resource management and job scheduling
 - ↪ Simple Linux Utility for Resource Management <https://slurm.schedmd.com/>
 - ↪ Handles submission, scheduling, execution, and monitoring of **jobs**
 - ↪ official [documentation](#), [tutorial](#), [FAQ](#)

- **User jobs** have the following key characteristics:
 - ↪ set of requested resources:
 - ✓ number of computing resources: **nodes** (including all their CPUs and cores) or **CPUs** (including all their cores) or **cores**
 - ✓ amount of **memory**: either per node or per CPU
 - ✓ **(wall)time** needed for the users tasks to complete their work
 - ↪ a requested node **partition** (job queue)
 - ↪ a requested **quality of service** (QoS) level which grants users specific accesses
 - ↪ a requested **account** for accounting purposes

Batch Scheduling Configuration (Slurm)

Slurm on ULHPC clusters

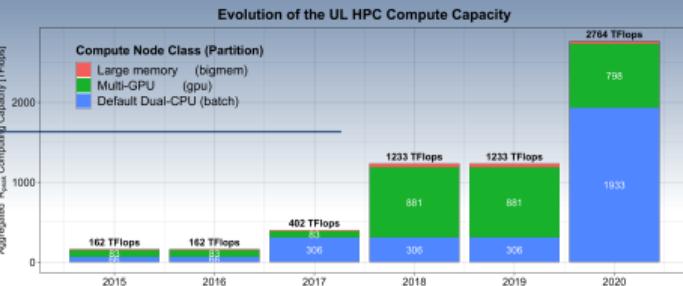
- Predefined **Queues/Partitions** depending on node type
 - ↳ batch (Default Dual-CPU nodes)
 - ↳ gpu (GPU nodes nodes)
 - ↳ bigmem (Large-Memory nodes)
 - ↳ In addition: interactive (for quicks tests)
 - ✓ for code development, testing, and debugging



- Max:** 64 nodes, 2 days walltime
Max: 4 nodes, 2 days walltime
Max: 1 node, 2 days walltime
Max: 2 nodes, 2h walltime

Slurm on ULHPC clusters

- Predefined **Queues/Partitions** depending on node type
 - ↳ batch (Default Dual-CPU nodes)
 - ↳ gpu (GPU nodes nodes)
 - ↳ bigmem (Large-Memory nodes)
 - ↳ In addition: interactive (for quicks tests)
 - ✓ for code development, testing, and debugging
- Queue Policy: **cross-partition QOS**, mainly tied to **priority level** (low → urgent)
 - ↳ long QOS with extended Max walltime (MaxWall) set to **14 days**
 - ↳ special **preemptible QOS** for best-effort jobs: **besteffort**.



Max: 64 nodes, 2 days walltime

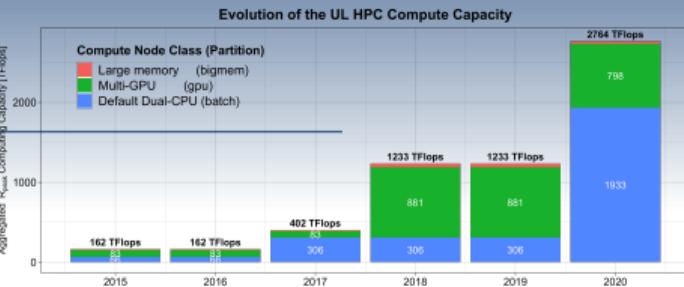
Max: 4 nodes, 2 days walltime

Max: 1 node, 2 days walltime

Max: 2 nodes, 2h walltime

Slurm on ULHPC clusters

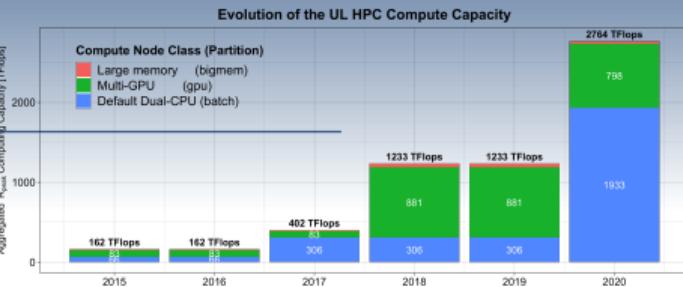
- Predefined **Queues/Partitions** depending on node type
 - ↪ batch (Default Dual-CPU nodes)
 - ↪ gpu (GPU nodes nodes)
 - ↪ bigmem (Large-Memory nodes)
 - ↪ In addition: interactive (for quicks tests)
 - ✓ for code development, testing, and debugging
- Queue Policy: **cross-partition QOS**, mainly tied to **priority level** (low → urgent)
 - ↪ long QOS with extended Max walltime (MaxWall) set to **14 days**
 - ↪ special **preemptible QOS** for best-effort jobs: **besteffort**.
- Accounts associated to supervisor (multiple associations possible)
 - ↪ Proper group/user accounting



Max: 64 nodes, 2 days walltime
Max: 4 nodes, 2 days walltime
Max: 1 node, 2 days walltime
Max: 2 nodes, 2h walltime

Slurm on ULHPC clusters

- Predefined **Queues/Partitions** depending on node type
 - ↪ batch (Default Dual-CPU nodes)
 - ↪ gpu (GPU nodes nodes)
 - ↪ bigmem (Large-Memory nodes)
 - ↪ In addition: interactive (for quicks tests)
 - ✓ for code development, testing, and debugging
- Queue Policy: **cross-partition QOS**, mainly tied to **priority level** (low → urgent)
 - ↪ long QOS with extended Max walltime (MaxWall) set to **14 days**
 - ↪ special **preemptible QOS** for best-effort jobs: **besteffort**.
- Accounts associated to supervisor (multiple associations possible)
 - ↪ Proper group/user accounting
- Slurm Multi-cluster configuration between **iris** and **aion**
 - ↪ easily **query** state on remote cluster {squeue...} -M, --cluster aion|iris



Max: 64 nodes, 2 days walltime

Max: 4 nodes, 2 days walltime

Max: 1 node, 2 days walltime

Max: 2 nodes, 2h walltime

Main Slurm Commands: Submit Jobs

```
$> sbatch -p <partition> [--qos <qos>] [-A <account>] [...] <path/to/launcher.sh>
```

Submitting Jobs

- **sbatch**: Submit batch **launcher script** for later execution **batch/passive mode**
 - ↪ allocate resources (nodes, tasks, partition, etc.)
 - ↪ runs a single **copy** of the batch script on the **first** allocated node

Main Slurm Commands: Submit Jobs

```
$> srun -p <partition> [--qos <qos>] [-A <account>] [...] --pty bash
```

Submitting Jobs

- **sbatch**: Submit batch **launcher script** for later execution **batch/passive mode**
 - ↪ allocate resources (nodes, tasks, partition, etc.)
 - ↪ runs a single **copy** of the batch script on the **first** allocated node
- **srun**: initiate parallel **job steps within a job OR start an interactive job**
 - ↪ allocate resources (number of nodes, tasks, partition, constraints, etc.)
 - ↪ launch a job that will execute on them.

Main Slurm Commands: Submit Jobs

```
$> salloc -p <partition> [--qos <qos>] [-A <account>] [...] <command>
```

Submitting Jobs

- **sbatch**: Submit batch **launcher script** for later execution **batch/passive mode**
 - ↪ allocate resources (nodes, tasks, partition, etc.)
 - ↪ runs a single **copy** of the batch script on the **first** allocated node
- **srun**: initiate parallel **job steps within a job OR start an interactive job**
 - ↪ allocate resources (number of nodes, tasks, partition, constraints, etc.)
 - ↪ launch a job that will execute on them.
- **salloc**: request interactive jobs/allocations
 - ↪ allocate resources (nodes, tasks, partition, etc.), either run a command or start a shell.

Specific Resource Allocation

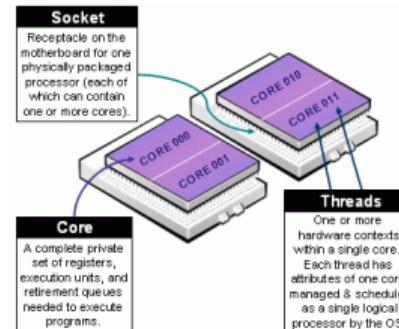
- Beware of Slurm terminology in Multicore Architecture!

→ Slurm Node = Physical node

✓ Advice: explicit number of expected tasks **per node**

-N <#nodes>

--ntasks-per-node <n>

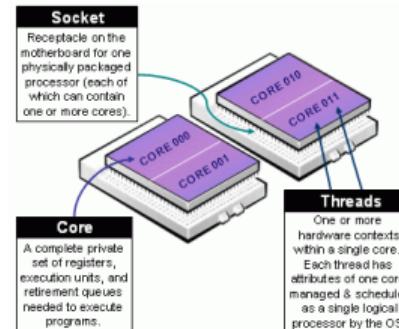


Specific Resource Allocation

- Beware of Slurm terminology in Multicore Architecture!

- Slurm Node = Physical node
 - ✓ **Advice:** explicit number of expected tasks **per node**
- Slurm Socket = Physical Socket/**CPU/Processor**

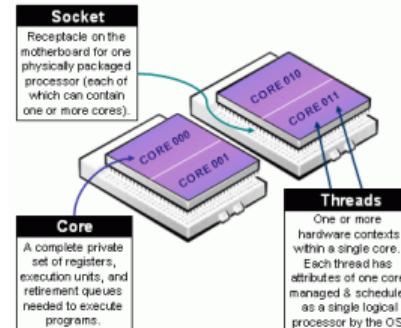
```
-N <#nodes>
--ntasks-per-node <n>
--ntasks-per-socket <n>
```



Specific Resource Allocation

- Beware of Slurm terminology in Multicore Architecture!

- Slurm Node = Physical node
 - ✓ Advice: explicit number of expected tasks **per node**
 - Slurm Socket = Physical Socket/CPU/Processor
 - Slurm CPU = Physical Core
 - ✓ Hyper-Threading (HT) Technology is **disabled** on all the compute nodes
 - ✓ #cores = #threads
 - ✓ Total number of tasks: \${SLURM_NTASKS}
- N <#nodes>
 --ntasks-per-node <n>
 --ntasks-per-socket <n>
 -c <#threads>
- $-c <N> \rightarrow OMP_NUM_THREADS=\${SLURM_CPUS_PER_TASK}$
 $\rightarrow srun -n \${SLURM_NTASKS} [...]$



Specific Resource Allocation

- Beware of Slurm terminology in Multicore Architecture!

- Slurm Node = Physical node
 - ✓ Advice: explicit number of expected tasks **per node**
- Slurm Socket = Physical Socket/CPU/Processor
- Slurm CPU = Physical Core
 - ✓ Hyper-Threading (HT) Technology is **disabled** on all the compute nodes
 - ✓ #cores = #threads -c <N> → OMP_NUM_THREADS=\${SLURM_CPUS_PER_TASK}
 - ✓ Total number of tasks: \${SLURM_NTASKS} → srun -n \${SLURM_NTASKS} [...]

- Important: Always align resource specs with physical NUMA characteristics

- Ex (AION): 16 cores per socket, 8 sockets (“physical” CPUs) per node (128c/node)
- [-N <N>] --ntasks-per-node <8n> --ntasks-per-socket <n> -c <thread>
 - ✓ Total: <N>×8×<n> tasks, each on <thread> threads
 - ✓ Ensure <n>×<thread>= 16 on aion
 - ✓ Ex: -N 2 --ntasks-per-node 32 --ntasks-per-socket 4 -c 4 (**Total:** 64 tasks)

Specific Resource Allocation

- Beware of Slurm terminology in Multicore Architecture!

- Slurm Node = Physical node
 - ✓ Advice: explicit number of expected tasks **per node**
- Slurm Socket = Physical Socket/CPU/Processor
- Slurm CPU = Physical Core
 - ✓ Hyper-Threading (HT) Technology is **disabled** on all the compute nodes
 - ✓ #cores = #threads -c <N> → OMP_NUM_THREADS=\${SLURM_CPUS_PER_TASK}
 - ✓ Total number of tasks: \${SLURM_NTASKS} → srun -n \${SLURM_NTASKS} [...]

- Important: Always align resource specs with physical NUMA characteristics

- Ex (IRIS): 14 cores per socket, 2 sockets (“physical” CPUs) per node (28c/node)
- [-N <N>] --ntasks-per-node <2n> --ntasks-per-socket <n> -c <thread>
 - ✓ Total: <N>×2×<n> tasks, each on <thread> threads
 - ✓ Ensure <n>×<thread>= 14 on iris
 - ✓ Ex: -N 2 --ntasks-per-node 4 --ntasks-per-socket 2 -c 7 (**Total:** 8 tasks)

Specific Resource Allocation

- Beware of Slurm terminology in Multicore Architecture!

- Slurm Node = Physical node
 - ✓ Advice: explicit number of expected tasks **per node**
 - Slurm Socket = Physical Socket/CPU/Processor
 - Slurm CPU = Physical Core
 - ✓ Hyper-Threading (HT) Technology is **disabled** on all the compute nodes
 - ✓ #cores = #threads
 - ✓ Total number of tasks: \${SLURM_NTASKS}
- N <#nodes>
--ntasks-per-node <n>
--ntasks-per-socket <n>
-c <#threads>
- c <N> → OMP_NUM_THREADS=\${SLURM_CPUS_PER_TASK}
→ srun -n \${SLURM_NTASKS} [...]

Hostname	Node type	#Nodes	#Socket	#Cores	RAM	Features
aion-[0001-0318]	Regular	318	8	128	256 GB	batch,epyc
iris-[001-108]	Regular	108	2	28	128 GB	batch,broadwell
iris-[109-168]	Regular	60	2	28	128 GB	batch,skylake
iris-[169-186]	Multi-GPU	18	2	28	768 GB	gpu,skylake,volta
iris-[191-196]	Multi-GPU	6	2	28	768 GB	gpu,skylake,volta32
iris-[187-190]	Large Memory	4	4	112	3072 GB	bigmem,skylake

- List available features: **sfeatures**

```
# OR: sinfo -o '%20N %.6D %.6c %15F %12P %f'
```

Uni.lu HPC School 2021/ PS3

Main Slurm Commands: Submit Jobs options

```
$> {sbatch | srun | salloc} [...]
```

Command-line option	Description	Example
-N <N> --ntasks-per-socket=<n> --ntasks-per-node=<n> -c <c> --mem=<m>GB -t [DD-]HH[:MM:]SS> -G <gpu> -C <feature>	<N> Nodes request <n> Tasks-per-socket request <n> Tasks-per-node request <c> Cores-per-task request (multithreading) <m>GB memory per node request Walltime request <gpu> GPU(s) request Feature request (Ex: broadwell,skylake,...)	-N 2 --ntasks-per-socket=14 --ntasks-per-node=28 -c 1 --mem 0 -t 4:00:00 -G 4 -C skylake
-p <partition> --qos <qos> -A <account>	Specify job partition/queue Specify job qos Specify account	
-J <name> -d <specification> --mail-user=<email> --mail-type=<type>	Job name Job dependency Specify email address Notify user by email when certain event types occur.	-J MyApp -d singleton --mail-type=END,FAIL

Main Slurm Commands: Collect Information

- Partition (queue) and node status
 - eventually filter on specific job state (**R**:running / **PD**:pending / **F**:failed / **PR**:preempted)

```
$> squeue [-u <user>] [-p <partition>] [--qos <qos>] [-t R|PD|F|PR]
```

Main Slurm Commands: Collect Information

- Partition (queue) and node status
 - eventually filter on specific job state (**R**:running / **PD**:pending / **F**:failed / **PR**:preempted)

```
$> squeue [-u <user>] [-p <partition>] [--qos <qos>] [-t R|PD|F|PR]
```

- Show partition status, summarized status (-s), problematic nodes (-R), reservations (-T)

```
$> sinfo [-p <partition>] {-s | -R | -T | ...}
```

Main Slurm Commands: Collect Information

- Partition (queue) and node status
 - eventually filter on specific job state (**R**:running / **PD**:pending / **F**:failed / **PR**:preempted)

```
$> squeue [-u <user>] [-p <partition>] [--qos <qos>] [-t R|PD|F|PR]
```

- Show partition status, summarized status (-s), problematic nodes (-R), reservations (-T)

```
$> sinfo [-p <partition>] {-s | -R | -T | ...}
```

- View job, partition, nodes, reservation status

```
$> scontrol show { job <jobid> | partition [<part>] | nodes <node>| reservation... }
```

Main Slurm Commands: Collect Information

Command	Description
sinfo	Report system status (nodes, partitions etc.)
squeue [-u \$(whoami)]	display jobs[steps] and their state
seff <jobid>	get efficiency metrics of past job
scancel <jobid>	cancel a job or set of jobs.
scontrol show [...]	view and/or update system, nodes, job, step, partition or reservation status
sstat	show status of running jobs.
sacct [-X] -j <jobid> [...]	display accounting information on jobs.
sprio	show factors that comprise a jobs scheduling priority
smap	graphically show information on jobs, nodes, partitions

```
### Get statistics on past job
slist <jobid>
# sacct [-X] -j <jobid> --format User,JobID,Jobname%30,partition,state,time,elapsed,MaxRSS, \
#                               MaxVMSize,nnodes,ncpus,nodelist,AveCPU,ConsumedEnergyRaw
#
# seff <jobid>
```

ULHPC Slurm Partitions 2.0

-p, -partition=<partition>

```
$> {srun|sbatch|salloc|sinfo|squeue...} -p <partition> [...]
```

AION Partition	Type	#Node	PriorityTier	DefaultTime	MaxTime	MaxNodes
interactive	floating	318	100	30min	2h	2
batch		318	1	2h	48h	64

IRIS Partition	Type	#Node	PriorityTier	DefaultTime	MaxTime	MaxNodes
interactive	floating	196	100	30min	2h	2
batch		168	1	2h	48h	64
gpu		24	1	2h	48h	4
bigmem		4	1	2h	48h	1

ULHPC Slurm QOS 2.0

--qos=<qos>

```
$> {srun|sbatch|salloc|sinfo|squeue...} [-p <partition>] --qos <qos> [...]
```

QOS	Partition	Allowed [L1] Account	Prio	GrpTRES	MaxTresPJ	MaxJobPU	Flags
besteffort	*	ALL	1			100	NoReserve
low	*	ALL (default for CRP/externals)	10			2	DenyOnLimit
normal	*	Default (UL,Projects,...)	100			50	DenyOnLimit
long	*	UL,Projects,etc.	100	node=6	node=2	1	DenyOnLimit,PartitionTimeLimit
debug	interactive	ALL	150	node=8		2	DenyOnLimit
high	*	(restricted) UL,Projects,Industry	200			10	DenyOnLimit
urgent	*	(restricted) UL,Projects,Industry	1000			100 ?	DenyOnLimit

- **Cross-partition QOS**, mainly tied to **priority level** (low → urgent)
 - ↪ Simpler names than before (i.e. no more qos- prefix)
 - ↪ special **preemptible QOS** for best-effort jobs: **besteffort**

Usage Example: Interactive Jobs

si,si-[gpu|bigmem]

```
# Simple interactive job - match feature name with target partition
srun -p interactive --qos debug -C {batch,gpu,bigmem} [...] --pty bash
# OR, **better**
si[-gpu|-bigmem] [...]
```

- *Floating* interactive partition allow to access all nodes
 - ↪ no more dedicated resources: optimize the cluster for passive jobs
 - ↪ no guarantee if partition is full **YET** backfilling and priority ensure first served

Node Type	Slurm command	Helper script
regular	srun -p interactive --qos debug -C batch [-C {broadwell,skylake}] [...] --pty bash	si [...]
gpu	srun -p interactive --qos debug -C gpu [-C volta[32]] -G 1 [...] --pty bash	si-gpu [...]
bigmem	srun -p interactive --qos debug -C bigmem [...] --pty bash	si-bigmem [...]

Usage Example: Interactive Jobs

si,si-[gpu|bigmem]

```
# Simple interactive job - match feature name with target partition
srun -p interactive --qos debug -C {batch,gpu,bigmem} [...] --pty bash
# OR, **better**
si[-gpu|-bigmem] [...]
```

```
# Simple interactive job on batch partition
(access)$> si
# srun -p interactive --qos debug -C batch --mem-per-cpu 4096 --pty bash
(node)$> echo $SLURM_NTASKS
1
(node)$> echo $SLURM_JOBID
2166371
$> sq
# squeue -u <login>
```

Usage Example: Regular Jobs

sbatch [...]

- Restricted to **Max 2 days / Max 64 nodes** default QOS induced by the job_submit.lua plugin
 - ↪ enforce precision of project/training account (**-A <account>**) **when appropriate** i.e. for
 - ✓ project name <project>
 - ✓ lecture/course name: <lecture>
 - ↪ **do NOT reserve gpu nodes (-p gpu) WITHOUT -G [1-4]**
 - ✓ Iris compute nodes in the gpu partition have 4xVolta V100 GPUs and 768GB RAM

Node Type	Slurm command
regular	sbatch [-A <project>] -p batch [--qos {high,urgent}] [-C {broadwell,skylake}] [...]
gpu	sbatch [-A <project>] -p gpu -G <N> [--qos {high,urgent}] [-C volta[32]] [...]
bigmem	sbatch [-A <project>] -p bigmem [--qos {high,urgent}] [...]

Usage Example: Long Jobs

--qos long

```
# Select target partition to bypass default walltime restrictions
sbatch -p {batch | gpu | bigmem} --qos long [...]
```

- Extended Max walltime (MaxWall) set to **14 days**

EuroHPC/PRACE Recommendations

↪ yet **restricted usage**: Max 6 nodes, Max 2 nodes per Job, Max **4** Jobs per User

Node Type	Slurm command
regular	sbatch [-A <project>] -p batch --qos long [-C {broadwell,skylake}] [...]
gpu	sbatch [-A <project>] -p gpu --qos long [-C volta[32]] -G 1 [...]
bigmem	sbatch [-A <project>] -p bigmem --qos long [...]

Usage Example: Best-effort Jobs

--qos besteffort

- Special **preemptible QOS** for **best-effort jobs**
 - ↪ *lowest* priority
 - ↪ **deleted if another non-besteffort job wants resources** where they are running.
 - ↪ useful to maximize use of the cluster for short-term jobs
 - ✓ several limits inherent to *regular* QOS are **bypassed** in besteffort
 - ↪ many scientific applications support internal state saving and restart **Checkpoint-restart**
 - ✓ System-level CR possible with **DMTCP: Distributed MultiThreaded CheckPointing**
 - ✓ see the official **DMTCP** launchers

```
sbatch -p {batch | gpu | bigmem} --qos besteffort [...]
```

Live Job Statistics

```
$> scontrol show job 2166371
JobId=2166371 JobName=bash
  UserId=<login>(<uid>) GroupId=clusterusers(666) MCS_label=N/A
  Priority=12741 Nice=0 Account=ulhpc QOS=debug JobState=RUNNING Reason=None
  [...]
  SubmitTime=2020-12-07T22:08:25 EligibleTime=2020-12-07T22:08:25
  StartTime=2020-12-07T22:08:25 EndTime=2020-12-07T22:38:25
  [...]
  WorkDir=/mnt/irisgpfs/users/<login>
```

Node/Job Statistics

```
$> sinfo
```

PARTITION	AVAIL	TIMELIMIT	NODES	STATE	NODELIST
interactive	up	4:00:00	196	idle	iris-[001-196]
batch*	up	2-00:00:00	5	mix	[...]
batch*	up	2-00:00:00	127	alloc	[...]
batch*	up	2-00:00:00	36	idle	[...]
gpu	up	2-00:00:00	4	resv	iris-[186,191-193]
gpu	up	2-00:00:00	19	alloc	[...]
gpu	up	2-00:00:00	1	idle	iris-185
bigmem	up	2-00:00:00	4	mix	iris-[187-190]

Node/Job Statistics

```
$> sinfo
```

PARTITION	AVAIL	TIMELIMIT	NODES	STATE	NODELIST
interactive	up	4:00:00	196	idle	iris-[001-196]
batch*	up	2-00:00:00	5	mix	[...]
batch*	up	2-00:00:00	127	alloc	[...]
batch*	up	2-00:00:00	36	idle	[...]
gpu	up	2-00:00:00	4	resv	iris-[186,191-193]
gpu	up	2-00:00:00	19	alloc	[...]
gpu	up	2-00:00:00	1	idle	iris-185
bigmem	up	2-00:00:00	4	mix	iris-[187-190]

```
$> slist <JOBID>
```

```
# sacct -j <JOBID> --format User,JobID,Jobname%30,partition,state,time,elapsed,\
```

```
# MaxRSS,MaxVMSize,nnodes,ncpus,nodelist,AveCPU,ConsumedEnergyRaw
```

```
# seff <JOBID>
```

Basic Slurm Launcher Examples

Prefer `--ntasks-per-node` over `-n` !

```
#!/bin/bash -l          # <--- DO NOT FORGET '-l'  
### Request a single task using one core on one node for 5 minutes in the batch queue  
#SBATCH -N 1  
#SBATCH --ntasks-per-node=1  
#SBATCH -c 1  
#SBATCH --time=0-00:05:00  
#SBATCH -p batch  
  
print_error_and_exit() { echo "***ERROR*** $*"; exit 1; }  
# Safeguard for NOT running this launcher on access/login nodes  
module purge || print_error_and_exit "No 'module' command"  
# List modules required for execution of the task  
module load <....>  
# [...]
```

Basic Slurm Launcher Examples

Prefer `--ntasks-per-node` over `-n` !

```
#!/bin/bash -l
### Request 6 single core tasks equally spread across two nodes for 3 hours
#SBATCH -N 2
#SBATCH --ntasks-per-node=3
#SBATCH -c 1
#SBATCH --time=0-03:00:00
#SBATCH -p batch
print_error_and_exit() { echo "***ERROR*** $*"; exit 1; }
module purge || print_error_and_exit "No 'module' command"
module load <...>
# [...]
```

Total: 6 tasks within 2 nodes
`sbatch <launcher>`

Total: 9 tasks within 3 nodes
`sbatch -N 3 <launcher>`

Basic Slurm Launcher Examples

Prefer --ntasks-per-node over -n !

```
#!/bin/bash -l
### Request as many tasks as cores available on a single node for 3 hours
#SBATCH -N 1
#SBATCH --ntasks-per-node=28 # On iris; for aion, use --ntasks-per-node=128
#SBATCH -c 1
#SBATCH --time=0-03:00:00
#SBATCH -p batch
print_error_and_exit() { echo "***ERROR*** $*"; exit 1; }
module purge || print_error_and_exit "No 'module' command"
module load <...>
# [...]
```

Total (iris): 28 tasks within 1 node
sbatch <launcher>

Total (iris): 56 tasks within 2 nodes
sbatch -N 2 <launcher>

Basic Slurm Launcher Examples (Large memory)

```
#!/bin/bash -l
### Request one sequential task requiring half the memory of a regular iris node for 1 day
#SBATCH -J MyLargeMemorySequentialJob      # Job name
#SBATCH --mail-user=Your.Email@Address.lu  # mail me ...
#SBATCH --mail-type=end,fail              # ... upon end or failure
#SBATCH -N 1
#SBATCH --ntasks-per-node=1
#SBATCH -c 1
#SBATCH --mem=64GB           # if above 112GB: consider bigmem partition (USE WITH CAUTION)
#SBATCH --time=1-00:00:00
#SBATCH -p batch             # if above 112GB: consider bigmem partition (USE WITH CAUTION)
print_error_and_exit() { echo "***ERROR*** $*"; exit 1; }
module purge || print_error_and_exit "No 'module' command"
module load <...>
# [...]
```

Basic Slurm Launcher Examples (GPU)

```
#!/bin/bash -l
### Request one GPU tasks for 4 hours - dedicate 1/4 of available cores for its management
#SBATCH -N 1
#SBATCH --ntasks-per-node=1
#SBATCH -c 7
#SBATCH -G 1
#SBATCH --time=00:04:00
#SBATCH -p gpu
print_error_and_exit() { echo "***ERROR*** $*"; exit 1; }
module purge || print_error_and_exit "No 'module' command"
# You probably want to use more recent gpu/CUDA-optimized software builds
module use /opt/apps/resif/iris/2019b/gpu/modules
module load <....>
# This should report a single GPU (over 4 available per gpu node)
nvidia-smi
# [...]
```

Slurm [Generic] Launchers 2.0

```
#!/bin/bash -l          # <--- DO NOT FORGET '-l'  
###SBATCH --job-name=<name>  
###SBATCH --dependency singleton  
###SBATCH -A <account>  
#SBATCH --time=0-01:00:00 # 1 hour  
#SBATCH --partition=batch # If gpu: set '-G <gpus>'  
#SBATCH -N 1             # Number of nodes  
#SBATCH --ntasks-per-node=2  
#SBATCH -c 1              # multithreading per task  
#SBATCH -o %x-%j.out      # <jobname>-<jobid>.out  
print_error_and_exit() { echo "****ERROR*** $*"; exit 1; }  
# Load ULHPC modules  
[ -f /etc/profile ] && source /etc/profile  
export OMP_NUM_THREADS=${SLURM_CPUS_PER_TASK:-1}  
module purge || print_error_and_exit "No 'module' command"  
module load <...>  
srun [-n $SLURM_NTASKS] [...]
```

• Best-Practices

- use /bin/bash -l on top
- set **reasonable time limits**
- set (short) job name
- specify account
- --exclusive allocation?
- **Avoid** Job arrays
- consider singleton pipelining
 - ✓ job dep. made easy
- GPU jobs (gpu partition)
 - ✓ Set #GPUs with -G <n>
- Use \$SCRATCH for large/temporary storage
- consider night jobs
 - ✓ --begin=20:00

Summary

1 Introduction

2 Getting Started on ULHPC

3 Batch Scheduling Configuration (Slurm)

Slurm commands

ULHPC Slurm Configuration

Usage Example

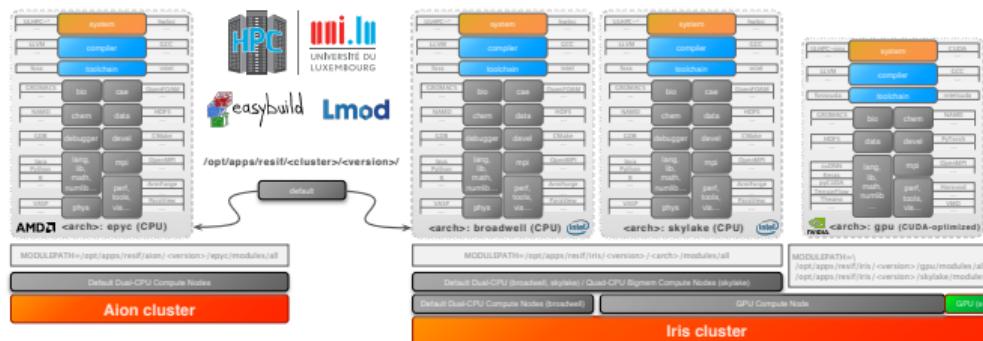
Basic Slurm Launchers

4 User [Software] Environment

Software/Modules Management

<https://hpc.uni.lu/user/software>

- Based on Environment Modules / LMod
 - ↪ convenient way to dynamically change the users environment with `module` command
- Currently on UL HPC: > 270 software packages, in *multiple* versions
 - ↪ software environment generated using RESIF 3.0 framework over Easybuild
 - ↪ hierarchical organization (Categorized Naming Scheme) Ex: `toolchain/{foss,intel}`



\$> `module load <category>/<software>[/<version>]`

Software/Modules Management

- Key module variable: \$MODULEPATH / where to look for modules.

↳ **default iris:**

/opt/apps/resif/iris/<version>/{{broadwell,skylake,gpu}}/modules/all

↳ **default aion:**

/opt/apps/resif/aion/<version>/{{epyc}}/modules/all

Software/Modules Management

- Key module variable: \$MODULEPATH / where to look for modules.

↳ **default iris:**

/opt/apps/resif/iris/<version>/{{broadwell,skylake,gpu}}/modules/all

↳ **default aion:**

/opt/apps/resif/aion/<version>/{{epyc}}/modules/all

Command	Description
module avail	Lists all the modules which are available to be loaded
module spider <pattern>	Search for among available modules (Lmod only)
module load <mod1> [mod2...]	Load a module
module unload <module>	Unload a module
module list	List loaded modules
module purge	Unload all modules (purge)
module use <path>	Prepend the directory to the MODULEPATH environment variable
module unuse <path>	Remove the directory from the MODULEPATH environment variable

ULHPC Toolchains and Software Set Versioning

- **Yearly** release based on Easybuild release of toolchains
 - see Component versions (**fixed per release**) in the **foss** and **intel** toolchains
 - ✓ count 6 months of validation/import after EB release before ULHPC release

Name	Type	2019b (legacy)	2020a	2020b (prod)	2021a (devel)
GCCCore	compiler	8.3.0	9.3.0	10.2.0	10.3.0
foss	toolchain	2019b	2020a	2020b	2021a
intel	toolchain	2019b	2020a	2020b	2021a
binutils		2.32	2.34	2.35	2.36
Python		3.7.4 (and 2.7.16)	3.8.2 (and 2.7.18)	3.8.6	3.9.2
LLVM	compiler	9.0.1	10.0.1	11.0.0	11.1.0
OpenMPI	MPI	3.1.4	4.0.3	4.0.5	4.1.1

```
# test (new) development software set
resif-load-swset-devel
# Restore production settings
resif-load-swset-prod
```

Interaction with Slurm and the software environment

Your Turn!

Hands-on Slurm

▶ url ◀ | [github](#) | [src](#)

- Run interactive job
- Run passive jobs
- Write basics launcher scripts

si
sbatch

Hands-on Modules

▶ url ◀ | [github](#) | [src](#)

- play with module command
- Use the MODULEPATH variable in order to change the software set
- Compile and run examples in at least two different languages

Thank you for your attention...

Questions?

ulhpc-tutorials.rtfd.io/en/latest/beginners



High Performance Computing @ Uni.lu

University of Luxembourg, Belval Campus
Maison du Nombre, 4th floor
2, avenue de l'Université
L-4365 Esch-sur-Alzette
mail: hpc@uni.lu

1 Introduction

2 Getting Started on ULHPC

3 Batch Scheduling Configuration (Slurm)

Slurm commands
ULHPC Slurm Configuration
Usage Example
Basic Slurm Launchers

4 User [Software] Environment

Uni.lu HPC School 2021 Contributors

	Dr. Xavier Besseron Research Scientist		Hyacinthe Cartiaux Infra. & HPC Arch. Engineer		Dr. Tiago C. Pessoa Postdoctoral Researcher
	Dr. Aurelien Ginohac Research Scientist		Dr. Emmanuel Kieffer Research Scientist		Sarah Peter Infra. & Arch. Engineer
	Dr. Loizos Koutsantonis Postdoctoral Researcher		Dr. Teddy Valette Infra. & HPC Arch. Engineer		Dr. Sébastien Varrette Research Scientist
	Dr. Ezhilmathi Krishnasamy Postdoctoral Researcher		Arlyne Vandeventer Project Manager	... and additional help (Survey, session tests)	

hpc.uni.lu